

DDPG and Soft AC

Spring, 2024

- 1 Deep Deterministic Policy Gradient
- 2 Soft Actor-Critic

Table of Contents

1 Deep Deterministic Policy Gradient

2 Soft Actor-Critic

Q-learning-DDPG-TD3-SAC

- ① Value based RL methods starting from Q-learning are also being developed over the years
- ② **DDPG**: Deterministic Policy Gradient Algorithms, Silver et al. ICML 2014
- ③ **TD3**: Addressing Function Approximation Error in Actor-Critic Methods, Fujimoto et al. ICML 2018
- ④ **SAC**: Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor, Haarnoja et al. ICML 2018

Stochastic Policy

- The policy is **stochastic** and denoted by

$$\pi_{\theta} : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A}),$$

where $\mathcal{P}(\mathcal{A})$ is the set of probability measures on \mathcal{A} and $\theta \in \mathbb{R}^n$ is a vector of n parameters.


- $\pi_{\theta}(a_t | s_t)$ is the conditional **probability density** at a_t associated with the policy.
- For the same state s , the actions stochastically selected according to θ might be different.

Deterministic Policy


- Select a **deterministic** action:

$$a = \mu_{\theta}(s)$$

- action is *uniquely* determined at the state s w.r.t the parameter θ
- suitable for the continuous action space, especially if the action space has many dimensions



Continuous Action Space




$$A \in [0, 2\pi] * [0, \pi]$$

DQN for Continuous Action Space



Deterministic Policy Gradient



Deep Deterministic Policy Gradient (DDPG)

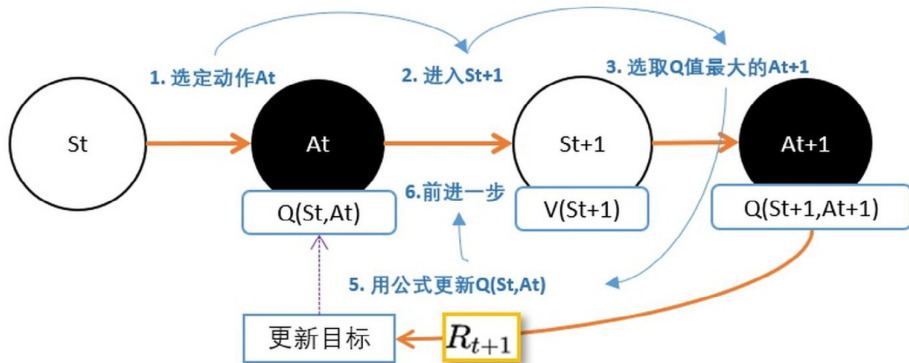
- 1 Motivation: how to extend DQN to the environment with continuous action space?
- 2 DDPG is very similar to DQN, which can be considered as a continuous action version of DQN

$$\text{DQN} : a^* = \arg \max_a Q^*(s, a)$$

$$\text{DDPG} : a^* = \arg \max_a Q^*(s, a) \approx Q_\phi(s, \mu_\theta(s))$$

- 1 a deterministic policy $\mu_\theta(s)$ directly gives the action that maximizes $Q_\phi(s, a)$
- 2 as action a is continuous we assume Q-function $Q_\phi(s, a)$ is differentiable with respect to a

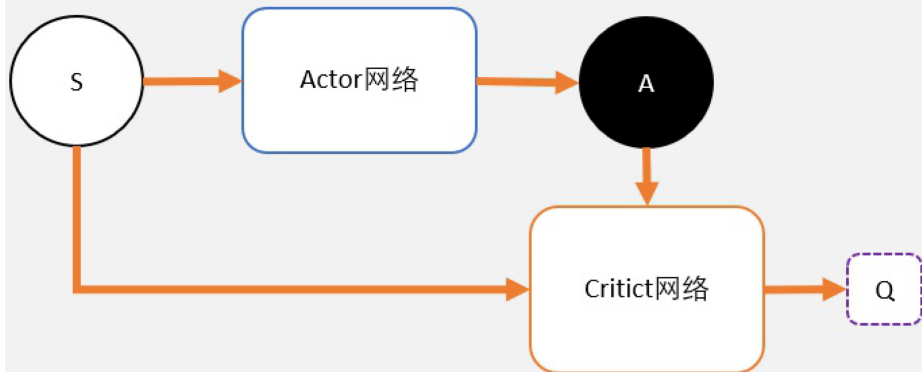
Q Function of DQN



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left[R + \gamma \max_a Q(S', a) - Q(S, A) \right]$$

Actor Network of DDPG

Actor的任务是找出A，令输出Q最大化



Actor Network of DDPG (2)

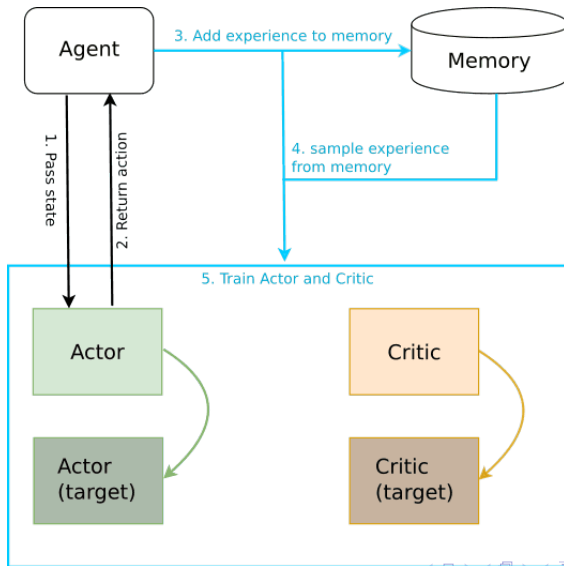
Critic 输出的价值代表了 Actor 预测动作的好坏, 因此策略网络的目标是最大化价值 *Value*, 自然就想到了用梯度上升法来最大化 $q(s, a; w)$, 于是, 我们可以对 $q(s, a; w)$ 求 θ 的梯度, 让我们将策略网络记作 $\pi(s; \theta)$:

$$pg = \frac{\partial q(s, \pi(s; \theta); w)}{\partial \theta} = \frac{\partial q(s, a; w)}{\partial a} * \frac{\partial a}{\partial \theta}$$

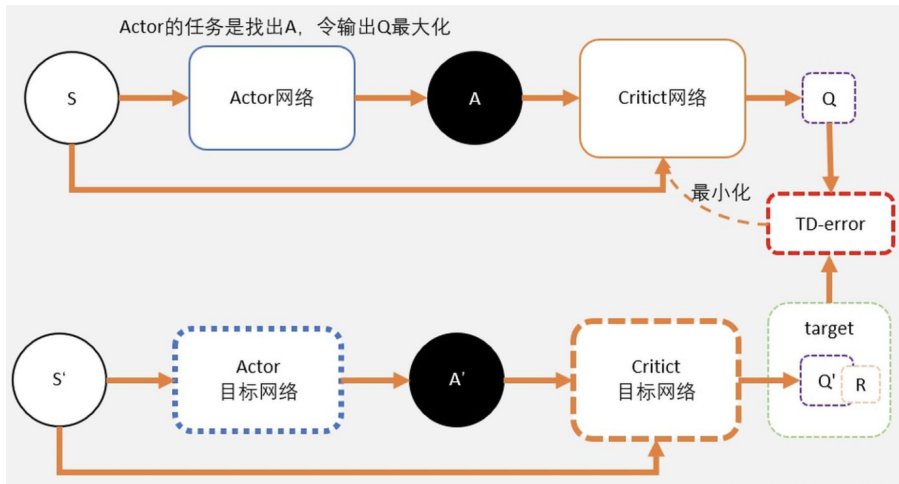
然后用梯度上升更新 θ :

$$\theta \leftarrow \theta + \alpha' * pg$$

DDPG Framework



DDPG Framework (2)



DDPG Algorithm

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for t = 1, T **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|\theta^{Q'}$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for

end for

DDPG Algorithm (2)

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M do

Initialize a random process \mathcal{N} for action exploration

Receive initial observation state s_1

for t = 1, T do

行动策略为随机策略

Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

Execute action a_t and observe reward r_t and observe new state s_{t+1}

Store transition (s_t, a_t, r_t, s_{t+1}) in R

经验回放

Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

目标网络

Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

Update the actor policy using the sampled gradient:

$$\nabla_{\theta^\mu} \mu|_{s_t} \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_t}$$

Update the target networks:

$$\begin{aligned} \theta^{Q'} &\leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'} \\ \theta^{\mu'} &\leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'} \end{aligned}$$

目标网络参数更新

end for
end for

Example: TORCS

`https://youtu.be/8CNck-hdys8`

Table of Contents

- 1 Deep Deterministic Policy Gradient
- 2 Soft Actor-Critic



Soft Actor-Critic (SAC)

- 1 SAC optimizes a stochastic policy in an off-policy way, which unifies stochastic policy optimization and DDPG-style approaches
- 2 SAC incorporates **entropy regularization**
- 3 Entropy is a quantity which measures how random a random variable is, $H(P) = E_{x \sim P}[-\log P(x)]$
- 4 Entropy-regularized RL: the policy is trained to maximize a trade-off between expected return and entropy, a measure of randomness in the policy

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{(s_t, a_t) \sim \rho_{\pi}} \left[\sum_t R(s_t, a_t) \right]$$

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{(s_t, a_t) \sim \rho_{\pi}} \left[\underbrace{\sum_t R(s_t, a_t)}_{\text{reward}} + \alpha \underbrace{H(\pi(\cdot | s_t))}_{\text{entropy}} \right]$$

Soft Actor-Critic (SAC)


 $-\log \pi(a'|s')$


$$q_{\pi}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') q_{\pi}(s', a')$$


$$q_{\pi}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') (q_{\pi}(s', a') - \alpha \log(\pi(a'|s')))$$

- The recursive Bellman equation for soft Q-function is

$$Q_{soft}(s_t, a_t) = r(s_t, a_t) + \gamma \mathbb{E}_{s_{t+1}, a_{t+1}} [Q_{soft}(s_{t+1}, a_{t+1}) - \alpha \log(\pi(a_{t+1}|s_{t+1}))]$$

Soft Actor-Critic (SAC)

Policy Iteration




Policy evaluation Estimate v_π

Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$

Greedy policy improvement



Soft Actor-Critic (SAC)

Policy iteration for traditional RL

- Policy evaluation: fix policy, update Q-function

$$Q_{\pi}(s, a) = r(s, a) + \lambda \mathbb{E}_{s', a'} Q_{\pi}(s', a')$$

- Policy improvement: update policy

$$\pi'(s) = \arg \max_a Q_{\pi}(s, a)$$

Policy iteration for SAC

- Policy evaluation: fix policy, update Q-function

$$Q_{soft}^{\pi}(s_t, a_t) = r(s_t, a_t) + \lambda \mathbb{E}_{s_{t+1}, a_{t+1}} [Q_{soft}^{\pi}(s_{t+1}, a_{t+1}) - \alpha \log(\pi(a_{t+1} | s_{t+1}))]$$

- Policy improvement: update policy

$$\pi' = \arg \min_{\pi_k \in \Pi} D_{KL}(\pi_k(\cdot | s_t) || \frac{\exp(\frac{1}{\alpha} Q_{soft}^{\pi}(s_t, \cdot))}{Z_{soft}^{\pi}(s_t)})$$